

Kaaval Assurance

**Routers predict.
Kaaval verifies.**

An assurance layer for AI decisions — checked locally on AMD compute before anyone downstream sees the answer.

VERIFY, NOT PREDICT. RECEIPTS, NOT PROMISES.

ONE PUBLIC CONTAINER

**AMD
+ Gemma
+ ROCm / vLLM**

Evidence Baseline — no credentials
Live Session — BYOK when needed

RECEIPTS, NOT PROMISES

A tribunal already rejected “the AI is separate.”

Moffatt v. Air Canada, BC Civil Resolution Tribunal, Feb 2024 — a real ruling, not a projection.

WHAT HAPPENED

Air Canada's chatbot told a customer he could book at full fare and claim a bereavement discount *after* travel. The real policy required the request *before* travel — the chatbot's answer was fluent, confident, and wrong.

Air Canada argued the chatbot was “a separate legal entity responsible for its own actions.” The tribunal rejected that outright and held the airline liable: \$812 CAD.


THE PATTERN KAAVAL GATES

The correct answer wasn't unknowable — it was sitting in a linked policy page in the same conversation. A fluent model still produced the wrong one.

This is exactly the shape of failure a deterministic contract check catches: not “is this text safe,” but “**does this claim match the policy on file.**”

Same shape of failure, caught before it ships.


A real contract from this repo: `support.refund_decision` — the \$500 cap is a range check, not a prompt.



Not a content filter. A guardrail blocks bad text. Kaaval decides which model answers — and proves what it cost.

The 5-step assurance path.

What each step does — and what breaks without it.



Only Layer 1 (Catch) gates acceptance. Adapt is Layer 2. A sampled offline audit (Layer 3) is display-only — it never gates a live answer.

EVIDENCE, SEPARATED HONESTLY

AMD proves execution. Fireworks proves selective economics.

Two artifacts, two claims — never blended into one result.

16 / 16

MEASURED

Local Contract-Conformance

Layer 1 · 16-case eval

324.6 / 479.6ms

MEASURED

Latency P50 / P95

16-case eval

14 / 16

MEASURED

Remote Calls Avoided

87.5% reduction

\$0.0333

CONFIGURED

Configured Cost Avoided

88.7% reduction

AMD/ATI host · 47.98 GiB VRAM · gfx1100 target · gemma-3-1b-it via ROCm + vLLM

Exact GPU marketing name unavailable; none inferred.

Configured-cost figure is a price estimate from recorded token counts, not a provider invoice — tagged CONFIGURED, not MEASURED, on purpose.

`bundle live-5be3acfa-amd-gemma-proof · checksummed`

WHY NOW

Enterprises deployed AI. They don't trust it unattended.

Every AI answer is becoming a business action — refunds, approvals, claims, severity calls. Adoption raced ahead of governance.

88%

of contact centers already run AI

adoption is not the question

25%

have fully integrated it unattended

the trust gap — Kaaval's wedge

\$15B

AI customer-service market, 2026

the surface Kaaval attaches to

BEACHHEAD

AI platform teams

Refund, triage, severity decisions where a wrong answer triggers a real action.

EXPAND

→ Regulated verticals

Support, telecom, finance, healthcare — every flow where “probably right” is a liability.

VISION

→ Every AI action, a receipt

The acceptance layer on top of inference spend — not the whole market, the governance slice of it.

Pull one public container. Evidence appears immediately.

Connect a model endpoint only when you want live assurance.

`ghcr.io/kaaval-ai/kaaval-assurance:act-ii`

linux/amd64 · publicly pullable · clean-smoke verified

- ✓ Evidence Baseline — measured AMD bundle, zero credentials
- ✓ Live Session — Fireworks BYOK or reachable Ollama / vLLM
- ✓ Safe defaults — no model download; paid remote and export closed
- ✓ Acceptance — 361 tests · linux/amd64 · UID 10001 (non-root)

NOW

Provider-neutral gateway

Observe decisions · Evaluate contracts · Produce receipts

30 DAYS

Validation before expansion

Blind semantic benchmark · Two-sided Layer 3 calibration · TCO experiment

DESIGN PARTNERS

Shadow-mode proof

Refunds and support · False accept / reject rates · Operational economics

Pull one container. Inspect measured AMD evidence. Connect your endpoint when you want live assurance.

Kaaval Assurance · Verify, don't predict. Receipts, not promises.