

The accountability layer for AI decisions.

When AI stops writing text and starts **taking actions** — refunds, approvals, claims — every action needs a record of whether it was allowed. Kaaval is that record.

VERIFY, DON'T PREDICT. · RECEIPTS, NOT PROMISES.

AI is moving from writing text to taking actions.

The last wave answered questions. This wave makes decisions — and a decision, unlike a sentence, has consequences the moment it ships.

YESTERDAY

AI wrote text

Summaries, drafts, chat. A wrong answer was an inconvenience — a human read it first.

TODAY

AI makes decisions

Approve a refund, route a claim, classify an incident, execute a tool call. The output IS the action.

THE GAP

No one is watching the decision

Monitoring shows latency and tokens. It cannot answer: did this decision cross the policy it was allowed to make?

A tribunal held a company liable for its AI's decision.

Moffatt v. Air Canada, 2024 — the chatbot invented a refund policy. The airline argued the AI was a separate entity. The tribunal rejected that.

WHAT HAPPENED

A fluent, confident, wrong answer triggered a real financial obligation before anyone checked it against policy. The company owned the outcome — and could not reconstruct why the decision was made.

THE PATTERN, EVERYWHERE

This is not an edge case — it is the default risk of every AI decision that acts before it is checked. The correct answer was knowable; the model produced the wrong one and it shipped.

That is the gap Kaaval closes: not “is this text safe,” but “did this decision match the policy on file — and can you prove it later.”

Three forces are converging on this exact moment.

01

Agents are taking real actions

Frameworks now let AI execute tools, move money, and change records autonomously. The blast radius of a wrong decision just went up.

02

Regulation makes it mandatory

The EU AI Act's high-risk record-keeping obligations take effect Aug 2, 2026. Penalties reach €15M or 3% of global turnover. Deployers stay liable even for vendor AI.

03

Adoption raced ahead of trust

88% of contact centers already run AI; only ~25% trust it unattended. Everyone deployed. Almost no one governs. That gap is the wedge.

WHAT KAAVAL IS

The acceptance layer between a model answer and a real action.

Every consequential AI decision gets three things it does not have today: a contract, a recovery path, and a replayable receipt.

A contract

An explicit, versioned rule for the task — checked in code, not asked in a prompt. The \$500 cap is a range check, not a suggestion.

A recovery path

A failed decision escalates once to a stronger model on the identical contract, or fails closed. It never ships nonconformant.

A receipt


Every decision is a redacted, replayable record — contract, checks, cost, provenance. Reconstructable later with zero model calls.

Not a content filter (that screens words). **Not a gateway** (that routes traffic). **Not an eval tool** (that tests offline). Kaaval governs the live decision and proves it.

HOW IT WORKS

Any model in. One governed decision out.

The request path is deterministic. No model is ever asked to judge another model.



✓ **conformant** ships ↻ **recovered** re-checked, ships ⚠ **no safe answer** withheld ⚡ **provider failure** not the model's fault

Implemented path: AssurancePipeline → Router → Verifier → TrajectoryStore. Layer-2 drift routing and a sampled, display-only Layer-3 audit sit around it.

Shipped, tested, and measured on real silicon.

16 / 16

MEASURED

contract-conformant on measured
AMD run

88.7%

CONFIGURED

fewer premium model calls (local-
first)

400+

MEASURED

tests · network-free · CI-clean

1

MEASURED

public container · SDK · terminal
cockpit

Gemma 3 on AMD ROCm + vLLM, gfx1100 · checksummed evidence bundle. Cost is a configured-price estimate from measured token counts — labeled configured, never presented as an invoice. We never blend measured and configured into one number. Every claim on this page maps to a stored field.

Zero-friction in. Never rip-and-replace.

Kaaval rides whatever you already run. Start in-process in twenty minutes; earn the request path.

TIER 0 · SDK

`@kaaval.assure`

One decorator on your decision function.
No new infra, no network hop. Shadow or enforce. Ships today.

SHIPPED

TIER 1 · SHADOW

Free diagnosis

Mirror your traffic. Two weeks later: how many decisions broke your own rules, and what local-first would save. Zero risk.

SHIPPED

TIER 2 · FILTER


On your gateway

An assurance filter on agentgateway / Envoy / LiteLLM — not a replacement. We govern the decision; they move the packets.

BUILDING

Assurance is the wedge. The record is the company.

Every guardian shares one join key — the model, the version, the decision. Together they become the accountability layer for all of agentic AI.



We claim only what ships. The platform is staged and honestly labeled — the vision is the destination, not a feature list we pretend is built.

The engine is copyable. What compounds is not.

A competitor can rebuild the verifier in a quarter. They cannot rebuild what only exists after you are in the decision path at volume.

The contract library

Validated, versioned rules per vertical — refunds, triage, claims, approvals. Fifty battle-tested contracts require fifty engagements to copy.

The evidence corpus

Cross-customer false-accept / false-reject rates per contract, per model, per version. “Which models can be trusted with refund decisions” only exists if you run at scale.

Audit system-of-record

When Kaaval receipts become a company's Article-12 evidence trail, ripping them out breaks audit continuity. Compliance formats are the stickiest lock-in in enterprise.

Provider-neutral by design

We ride every gateway and model; we're owned by none. As the connectivity layer commoditizes, the assurance layer on top is where the value concentrates.

Open-source adoption. Metered on assured decisions.

Land free and local; monetize when local receipts must become shared organizational memory.

OPEN SOURCE

Developer

- Local assurance SDK + contract authoring
- Local Kaaval Top cockpit
- Single-node receipts and replay

Free

TEAM

Hosted control plane

- Shared contracts, review queue, history
- Scheduled evidence + dashboards
- Metered per assured decision

Usage-based

ENTERPRISE

Self-hosted / VPC

- SSO, RBAC, retention, SIEM export
- Signed Article-12 evidence bundles
- Validated domain contract packs

Committed

Pricing is a hypothesis until design partners commit. The novel unit — price per assured decision — aligns cost with the value delivered.

Everyone else answers a different question.

Guardrails / Bedrock

“Is this content safe?”

One call. Screens words, not business rules. Can't prove a decision.

Patronus / evals

“Is the model good in general?”

Offline, pre-deployment. Doesn't govern the live decision that ships.

agentgateway (LF)

“Can these agents connect?”

Connectivity, routing, auth. No notion of whether a decision was allowed.

Datadog / Langfuse

“What happened?”

Observability after the fact. A trace, not an acceptance record.

Kaaval: “Was this decision allowed — and can you prove it?” *If you can't audit it, it's not governance. It's hope.*

Built by two engineers who refuse to overclaim.

Milind Gunjan

Product, systems, go-to-market

Hari Krishna Govindarajan

Architecture, assurance engine, platform

The vision is simple:

Every enterprise AI action should carry an acceptance receipt.

VERIFY, DON'T PREDICT. · RECEIPTS, NOT PROMISES.