The accountability layer for AI decisions.

When AI stops writing text and starts **taking actions** — refunds, approvals, claims — every action needs a record of whether it was allowed. Kaaval is that record.

VERIFY, DON'T PREDICT. · RECEIPTS, NOT PROMISES.

OUR MISSION

Make every AI decision accountable.

Autonomous AI will run more of the economy every year — approving, deciding, and acting with no human in the loop. It can only be trusted with that power if every decision it makes is **provable, recoverable, and on the record.** We are building that record. Kaaval is the trust layer for the agentic economy.

Every system society learned to trust got an accountability primitive first — **commerce got the ledger, the web got the certificate.** **Agentic AI gets the receipt.**

THE SHIFT

AI is moving from writing text to taking actions.

The last wave answered questions. This wave makes decisions — and a decision, unlike a sentence, has consequences the moment it ships.

YESTERDAY

AI wrote text

Summaries, drafts, chat. A wrong answer was an inconvenience — a human read it first.

→

TODAY

AI makes decisions

Approve a refund, route a claim, classify an incident, execute a tool call. The output IS the action.

→

THE GAP

No one is watching the decision

Monitoring shows latency and tokens. It cannot answer: did this decision cross the policy it was allowed to make?

THE PROBLEM IS ALREADY LITIGATED

A tribunal held a company liable for its AI's decision.

Moffatt v. Air Canada, 2024 — the chatbot invented a refund policy. The airline argued the AI was a separate entity. The tribunal rejected that.

WHAT HAPPENED

A fluent, confident, wrong answer triggered a real financial obligation before anyone checked it against policy. The company owned the outcome — and could not reconstruct why the decision was made.

THE PATTERN, EVERYWHERE

This is not an edge case — it is the default risk of every AI decision that acts before it is checked. The correct answer was knowable; the model produced the wrong one and it shipped.

That is the gap Kaaval closes: not “is this text safe,” but “did this decision match the policy on file — and can you prove it later.”

Source: BC Civil Resolution Tribunal, Moffatt v. Air Canada (Feb 2024) · publicly reported

WHY NOW

Three forces are converging on this exact moment.

01

Agents are taking real actions

Frameworks now let AI execute tools, move money, and change records autonomously. The blast radius of a wrong decision just went up.

02

Regulation makes it mandatory

The EU AI Act's high-risk record-keeping obligations take effect Aug 2, 2026. Penalties reach €15M or 3% of global turnover. Deployers stay liable even for vendor AI.

03

Adoption raced ahead of trust

88% of contact centers already run AI; only ~25% trust it unattended. Everyone deployed. Almost no one governs.

Gartner named the category: **“Guardian Agents” — runtime enforcement for agentic AI (Feb 2026). AI TRiSM: $3.1B → $13.8B by 2030 · 95% of C-suites report AI incidents · ungoverned incidents average $4.4M.**

Sources: EU AI Act Art. 12/19 (artificialintelligenceact.eu) · Gartner AI TRiSM Market Guide + Guardian Agents Market Guide (Feb 2026) · market reports 2026

WHAT KAAVAL IS

The acceptance layer between a model answer and a real action.

Every consequential AI decision gets three things it does not have today: a contract, a recovery path, and a replayable receipt.

A contract

An explicit, versioned rule for the task — checked in code, not asked in a prompt. The $500 cap is a range check, not a suggestion.

A recovery path

A failed decision escalates once to a stronger model on the identical contract, or fails closed. It never ships nonconformant.

A receipt

Every decision is a redacted, replayable record — contract, checks, cost, provenance. Reconstructable later with zero model calls.

Not a content filter (that screens words). **Not a gateway** (that routes traffic). **Not an eval tool** (that tests offline). Kaaval governs the live decision and proves it.

HOW IT WORKS

Any model in. One governed decision out.

The request path is deterministic. No model is ever asked to judge another model.

ANY MODEL	→	CONTRACT CHECK	→	FOUR OUTCOMES	→	RECEIPT
GPT-5.6 · Gemma · Claude · local		deterministic, in code · fails closed		conform · recover · withhold · provider-fail		replayable · redacted · audit-ready

✓ **conformant** ships ↻ **recovered** re-checked, ships ⊘ **no safe answer** withheld ⚡ **provider failure** not the model's fault

Implemented path: AssurancePipeline → Router → Verifier → TrajectoryStore. Layer-2 drift routing and a sampled, display-only Layer-3 audit sit around it.

A 50-year-old engineering pattern, applied to AI decisions: ***the Simplex runtime-assurance architecture — unverified controller, verified safety monitor, deterministic fallback.***

THIS IS BUILT, NOT A DECK

Shipped, tested, and measured on real silicon.

16 / 16

MEASURED

contract-conformant on measured AMD run

88.7%

CONFIGURED

fewer premium model calls (local-first)

400+

MEASURED

tests · network-free · CI-clean

1

MEASURED

public container · SDK · terminal cockpit

Gemma 3 on AMD ROCm + vLLM, gfx1100 · checksummed evidence bundle. Cost is a configured-price estimate from measured token counts — labeled configured, never presented as an invoice. We never blend measured and configured into one number. Every claim on this page maps to a stored field.

Evidence: docs/feature-verification.md · artifacts/demo-live- · fireworks-cost-comparison-*.json*

HOW YOU ADOPT IT

Zero-friction in. Never rip-and-replace.

Kaaval rides whatever you already run. Start in-process in twenty minutes; earn the request path.

TIER 0 · SDK

`@kaaval.assure`

One decorator on your decision function. No new infra, no network hop. Shadow or enforce. Ships today.

SHIPPED

TIER 1 · SHADOW

Free diagnosis

Mirror your traffic. Two weeks later: how many decisions broke your own rules, and what local-first would save. Zero risk.

SHIPPED

TIER 2 · FILTER

On your gateway

An assurance filter on agentgateway / Envoy / LiteLLM — not a replacement. We govern the decision; they move the packets.

BUILDING

WHERE THIS GOES

Assurance is the wedge. The record is the company.

Every guardian shares one join key — the model, the version, the decision. Together they become the accountability layer for all of agentic AI.

Assurance

Governs each decision at runtime. The wedge.

SHIPPED

NanoCanary

Measures how a model drifts under pressure, before release.

BUILDING

Substrate

Governed memory — what the agent knew, and what it ignored.

BUILDING

System of record

One replayable evidence chain for every AI action.

ROADMAP

We claim only what ships. The platform is staged and honestly labeled — the vision is the destination, not a feature list we pretend is built.

WHY THIS IS A COMPANY, NOT A FEATURE

The engine is copyable. What compounds is not.

A competitor can rebuild the verifier in a quarter. They cannot rebuild what only exists after you are in the decision path at volume.

The contract library

Validated, versioned rules per vertical — refunds, triage, claims, approvals. Fifty battle-tested contracts require fifty engagements to copy.

The evidence corpus

Cross-customer false-accept / false-reject rates per contract, per model, per version. “Which models can be trusted with refund decisions” only exists if you run at scale.

Audit system-of-record

When Kaaval receipts become a company's Article-12 evidence trail, ripping them out breaks audit continuity. Compliance formats are the stickiest lock-in in enterprise.

Provider-neutral by design

We ride every gateway and model; we're owned by none. As the connectivity layer commoditizes, the assurance layer on top is where the value concentrates.

BUSINESS MODEL

Open-source adoption. Metered on assured decisions.

Land free and local; monetize when local receipts must become shared organizational memory.

OPEN SOURCE

Developer

- Local assurance SDK + contract authoring
- Local Kaaval Top cockpit
- Single-node receipts and replay

Free

TEAM

Hosted control plane

- Shared contracts, review queue, history
- Scheduled evidence + dashboards
- Metered per assured decision

Usage-based

ENTERPRISE

Self-hosted / VPC

- SSO, RBAC, retention, SIEM export
- Signed Article-12 evidence bundles
- Validated domain contract packs

Committed

Pricing is a hypothesis until design partners commit. The novel unit — price per assured decision — aligns cost with the value delivered.

Everyone else answers a different question.

Guardrails / Bedrock	*"Is this content safe?"*	One call. Screens words, not business rules. Can't prove a decision.
Patronus / evals	*"Is the model good in general?"*	Offline, pre-deployment. Doesn't govern the live decision that ships.
agentgateway (LF)	*"Can these agents connect?"*	Connectivity, routing, auth. No notion of whether a decision was allowed.
Datadog / Langfuse	*"What happened?"*	Observability after the fact. A trace, not an acceptance record.

Kaaval: "Was this decision allowed — and can you prove it?" *If you can't audit it, it's not governance. It's hope.*

THE WORLD WE'RE BUILDING

When AI can be trusted with decisions, it can run real things.

Five years out, the accountability we're building isn't a product a few teams buy — it's the default every consequential AI action carries.

Billions of decisions a day

Every consequential AI action — a refund, a diagnosis, a trade, a dispatch — carries a verifiable record of what it was allowed to do and why.

The standard evidence format

Regulators, auditors, and insurers read Kaaval receipts the way finance reads a ledger. Proving an AI decision becomes a query, not an archaeology dig.

AI safe enough to run the economy

Autonomous agents operate money, healthcare, and infrastructure — because their decisions are accountable by default, not trusted by hope.

The agentic economy needs a trust layer. We intend to be it. *Stated as vision — the destination we build toward, not shipped features.*

TEAM & VISION

Built by two engineers who refuse to overclaim.

Milind Gunjan

Product, systems, go-to-market

Hari Krishna Govindarajan

Architecture, assurance engine, platform

The vision is simple:

Every enterprise AI action should carry an acceptance receipt.

VERIFY, DON'T PREDICT. · RECEIPTS, NOT PROMISES.